

# oneAPI Data Parallel C++ Library

For the DPC++ Technical Advisory Board discussion

April 2021

# Agenda

1. oneAPI Data Parallel C++ Library (oneDPL) recap
2. Notable oneDPL changes
3. Experimental Range-based API
4. Experimental Asynchronous API
5. Open questions

# oneAPI Data Parallel C++ Library: recap

- Complements Data Parallel C++ (DPC++) with the functionality similar to the standard C++ library
  - Extended set of STL algorithms with standard-like & special execution policies
  - “Extended subset” of standard C++ API supported in DPC++ kernels
- Goal is greater productivity and portability (comparing to “pure DPC++”)
- Specification:  
<https://spec.oneapi.com/versions/latest/elements/oneDPL/source/index.html>
- Open source implementation: <https://github.com/oneapi-src/oneDPL>
  - Based on LLVM PSTL project
  - Contributions are welcome!

# News since the last year

# Notable oneDPL changes

- The namespace was chosen: `oneapi::dpl`
  - Also, `::dpl` is a short alias
  - Dropped the idea of namespace `oneapi::std` for standard compliant API because of usability problems
- Algorithms are blocking by default
  - Intel's implementation provides a macro to allow non-blocking invocation
- Execution policy modifications
  - Class names: `device_policy`, `fpga_policy` (experimental)
  - Predefined objects: `dpcpp_default`, `dpcpp_fpga`
  - Implicit conversion to `sycl::queue`
- Multiple other improvements & adjustments

# Notable implementation-specific additions

- <random> implementation suitable for DPC++ kernels
- Experimental range-based API for some algorithms
- Experimental asynchronous API for some algorithms
- All these APIs are not yet a part of the oneDPL specification
  - Proposals to add these to the version 1.1 will be created at the [oneapi-src GitHub](#)

# <oneapi/dpl/random>

- Provides a subset of C++ <random>
- Standard compliant APIs with extensions to generate several RNs into `sycl::vec`
- Can be used on the host and in kernels

- Engine classes

<code>linear_congruential_engine</code>	<code>discard_block_engine</code>
<code>subtract_with_carry_engine</code>	<code>default_engine</code>

- Predefined engines

<code>minstd_rand[_vec]</code>	<code>minstd_rand0[_vec]</code>
<code>ranlux24_base[_vec]</code>	<code>ranlux48_base[_vec]</code>
<code>ranlux24[_vec]</code>	<code>ranlux48[_vec]</code>

- Distributions

<code>uniform_int_distribution</code>
<code>uniform_real_distribution</code>
<code>normal_distribution</code>

# <oneapi/dpl/random> usage example

```
template<int VecSize>
void random_fill(float* usmptr, std::size_t n) {

    auto zero = oneapi::dpl::counting_iterator<std::size_t>(0);

    std::for_each(oneapi::dpl::execution::dpcpp_default,
                  zero, zero + n/VecSize,
                  [usmptr](std::size_t i){

                        auto offset = i * VecSize;

                        oneapi::dpl::minstd_rand_vec<VecSize> engine(seed, offset);
                        oneapi::dpl::uniform_real_distribution<sycl::vec<float, VecSize>> distr;

                        auto res = distr(engine);
                        res.store(i, sycl::global_ptr<float>(usmptr));

                    });
}
```



# Experimental range-based API

# Range-based API for algorithms

- C++20 adds Ranges into the C++ standard library
  - Very powerful and expressive functional API
  - But does not yet support execution policies
- We work on adding range support for oneDPL algorithms
  - Only for a subset of the standard algorithms and views
  - Not fully standard-compliant (not based on concepts, no projections, ...)
- Available as an experimental feature since oneDPL v2021.1

# Ranges: programmability and performance

## A pipeline of 3 kernels:

```
std::reverse(policy, begin(data), end(data));  
std::transform(policy, begin(data), end(data), begin(result), [](auto i){return i*i;});  
auto res = std::find_if(policy, begin(result), end(result), pred);
```

## With fancy iterators (1 kernel):

```
auto res = std::find_if(policy,  
    make_transform_iterator(make_reverse_iterator(end(data)), [](auto i){return i*i;}),  
    make_transform_iterator(make_reverse_iterator(begin(data)), [](auto i){return i*i;}),  
    pred);
```

## With ranges (1 kernel):

```
using namespace oneapi::dpl::experimental::ranges;  
auto res = find_if(policy,  
    data | views::reverse | views::transform([](auto i){return i*i;}),  
    pred);
```

And it can be quite faster than the original 3-kernel variant

# Using ranges with DPC++ execution policies

- Make pipelines based on a factory
  - *iota* factory generates a sequence of indexes on the fly:  
`auto view = views::iota(0, n) | views::transform(lambda);`
- Make pipelines transforming data in USM
  - *subrange* creates a range view over two USM pointers  
`auto view = views::subrange(ptr, ptr+n) | views::transform(lambda);`
- Make pipelines transforming data in a SYCL buffer
  - *views::all* creates a range view over SYCL buffer  
`auto view = views::all(buf) | views::transform(lambda);`
  - Variations of it to specify the type of access: `all_read`, `all_write`
  - The data access operations may only be executed on device
- **Requirement:** ranges, views, pipelines of views must be **device copyable** to use them with DPC++ execution policies
  - How can we ensure that for the standard library implementations of GCC & LLVM?

# Range-based API (as of oneDPL v2021.3)

## Algorithms (34)

any_of	sort
all_of	stable_sort
for_each	is_sorted_until
find_if	is_sorted
find_if_not	equal
find	move
find_end	merge
find_first_of	min_element
count	max_element
count_if	minmax_element
search	reduce
search_n	transform_reduce
copy	exclusive_scan
transform	inclusive_scan
remove_if	transform_exclusive_scan
remove	transform_inclusive_scan
replace_if	
replace	

## Views

all  
all\_read  
all\_write  
subrange  
  
iota  
fill  
generate  
  
reverse  
transform  
rotate  
take  
drop  
zip  
permutation

# Experimental asynchronous API

# Blocking vs. asynchronous

- oneDPL algorithms with DPC++ execution policies are blocking
  - Standard-compliant: return when execution completes (on the device)
  - In some cases, may transfer data back to the host
- “Deferred waiting” mode can be enabled by a macro
  - Functions with no return value may submit a kernel and return immediately
  - Requires explicit waiting on the queue or buffer access
  - Non-standard, and not even in the oneDPL specification
- New: experimental explicitly asynchronous APIs
  - Functions never wait, instead returning a future-like object
  - Allows simultaneous use of multiple devices
  - Better exploits DPC++ asynchronous capabilities (hiding latencies etc.)
  - Enables composing oneDPL algorithms into static data flow graphs

# Looking at the API

- Adds \*\_async suffix to algorithm names
  - `oneapi::dpl::experimental::for_each_async(...);`
  - **Possible alternatives:** a different namespace, a different policy class
- May take an arbitrary number of dependencies as extra arguments
  - `for_each_async(policy, first, last, lambda [,events...]);`
- Returns an unspecified future-like type
  - `auto res = reduce_async(policy, first, last, std::plus<>{});`  
`auto sum = res.get();`
  - Also holds internal buffers; keep track of its lifetime!
  - Interoperable with “native” backend specific types, e.g. `sycl::event`
  - **Possible alternatives:** a generic future type or backend specific futures



# Try it: <oneapi/dpl/async>

- Algorithms supported since oneDPL v2021.2
  - copy\_async, fill\_async, for\_each\_async, reduce\_async, transform\_async, transform\_reduce\_async, sort\_async
- More to be added
- Your feedback is welcome!

```
int async_example(sycl::buffer<int>& a) {  
    auto& pol = dpl::execution::dpcpp_default;  
  
    /* Build and compute a simple dependency chain:  
       Fill buffer -> Transform -> Reduce */  
    auto fut1 = dpl::experimental::fill_async(  
        pol, dpl::begin(a), dpl::end(a), 7);  
  
    auto fut2 = dpl::experimental::transform_async(  
        pol, dpl::begin(a), dpl::end(a),  
        dpl::begin(a), [](int x){return x+1;},  
        fut1);  
  
    return dpl::experimental::reduce_async(  
        pol, dpl::begin(a), dpl::end(a),  
        fut2).get(); // blocks to get the value  
}
```

# Other open questions

# The minimal supported version of C++

Currently, oneDPL claims support for C++11 but:

- SYCL 2020 requires C++17
- DPC++ compiler does not support C++11 well
- Range-based API implementation requires at least C++17
  - Also the standard library of GCC 8.1 / LLVM 7 or higher
  - De-jure, ranges are a C++20 feature
- There is a strong desire to switch to C++17 (or at least C++14)
- Do you think it's acceptable for the specification and/or the implementation?

# Policy renaming/rebinding

- Sometimes one may want to name DPC++ kernels
  - For debugging, profiling, or lack of support for unnamed ones
- With oneDPL, *unique* names need to be given to each policy object

```
dpl::execution::device_policy<MyKernelName> pol;
```
- A policy can be “renamed” at algorithm invocation

```
for_each( dpl::execution::make_device_policy<NewKernelName>(pol), ... );
```

  - Creates a temporary object used only once.
  - But the syntax looks overly verbose and non-obvious
- Looking for something simpler, e.g.

```
for_each( dpl::execution::rename<NewKernelName>(pol), ... );  
for_each( pol.template rebind<NewKernelName>(), ... );
```
- Opinions?

# Closing

# Call to action

- Try out the experimental API: ranges and asynchronous algorithms  
<https://github.com/oneapi-src/oneDPL>
  - Provide your comments & feedback on the API
  - Participate in review / discussions of specification update proposals  
<https://github.com/oneapi-src/oneAPI-spec/discussions>
  - Share information with others who might be interested to contribute
- 
- Thanks!

# Backup

Materials from 2020 reviews

# A usage example for ranges

- A code sample (derived from SyclParallelSTL project)

```
// compute the number characters that start a new word
int wc = std::inner_product(                // what?? Let's figure out
    first, last,                          // a sequence of characters
    first + 1,                            // same sequence shifted by 1 character
    0,                                    // initialize the sum to 0
    std::plus<int>{},                     // combine partial sums
    [](char s1, char s2) {                // check two consequent chars to find where a word starts
        return is_word_start(s1, s2)? 1: 0;
    });
```

- The same thing done with ranges:

```
// compute the number characters that start a new word
using namespace std::ranges::views;
int wc = std::ranges::count_if(             // the algorithm name is descriptive
    zip( text|take(text.size()-1), text|drop(1) ), // views are easy to combine
    [](auto v) {                             // the predicate gets a tuple created by the zip view
        return is_word_start(std::get<0>(v), std::get<1>(v)); // working with tuples is cumbersome
    });
```